

Name: Abu Butt

Professor: Izidor Gertner

Class: Parallel Processing

Date: July 19, 2017

Topic: Mandelbrot Set

Introduction:

The Mandelbrot set is used to refer both to a general class of fractal sets and to a particular instance of such a set. In general, a Mandelbrot set marks the set of points in the complex plane such that the corresponding Julia set is connected and not computable. The Mandelbrot set is obtained from the quadratic recurrence equation $z_{n+1} = z_n^2 + C$. with $z_0 = C$, where points C in the complex plane for which the orbit of z_n does not tend to infinity are in the set. Mandelbrot set images may be created by sampling the complex numbers and determining for each sample point c , where the result of iterating the above function goes to infinity.

Mandelbrot Psuedocode:

```
For each pixel (Px, Py) on the screen, do:
{
    x0 = scaled x coordinate of pixel (scaled to lie in the Mandelbrot X
scale (-2.5, 1))
    y0 = scaled y coordinate of pixel (scaled to lie in the Mandelbrot Y
scale (-1, 1))
    x = 0.0
    y = 0.0
    iteration = 0
    max_iteration = 1000
    // Here N=2^8 is chosen as a reasonable bailout radius.
    while ( x*x + y*y < (1 << 16) AND iteration < max_iteration ) {
        xtemp = x*x - y*y + x0
        y = 2*x*y + y0
        x = xtemp
        iteration = iteration + 1
    }
    // Used to avoid floating point issues with points inside the set.
    if ( iteration < max_iteration ) {
        // sqrt of inner term removed using log simplification rules.
        log_zn = log( x*x + y*y ) / 2
        nu = log( log_zn / log(2) ) / log(2)
        // Rearranging the potential function.
        // Dividing log_zn by log(2) instead of log(N = 1<<8)
        // because we want the entire palette to range from the
        // center to radius 2, NOT our bailout radius.
        iteration = iteration + 1 - nu
    }
}
```

```
color1 = palette[floor(iteration)]
color2 = palette[floor(iteration) + 1]
// iteration % 1 = fractional part of iteration.
color = linear_interpolate(color1, color2, iteration % 1)
plot(Px, Py, color)
}
```

Above image is the pseudocode for the Mandelbrot set. Following the Mandelbrot pseudocode, we were to provide either C++ or C code and produce an image and count the timing using query performance using visual studio.

Mandelbrot Code in C++:

```

1  #include "SFML/Graphics.hpp"
2  #include<iostream>
3  #include<windows.h>
4  #include <tchar.h>
5  using namespace std;
6  //resolution of the window
7  const int width = 1280;
8  const int height = 720;
9
10 //used for complex numbers
11 struct complex_number
12 {
13     long double real;
14     long double imaginary;
15 };
16
17 void generate_mandelbrot_set(sf::VertexArray& vertexarray, int pixel_shift_x, int pixel_shift_y, int precision, float zoom)
18 //void generate_mandelbrot_set(sf::VertexArray& vertexarray, int pixel_shift_x, int pixel_shift_y, int precision)
19 {
20     #pragma omp parallel for
21     for (int i = 0; i < height; i++)
22     {
23         for (int j = 0; j < width; j++)
24         {
25             //scale the pixel location to the complex plane for calculations
26             //long double x = ((long double)j - pixel_shift_x);
27             //long double y = ((long double)i - pixel_shift_y);
28             long double x = ((long double)j - pixel_shift_x) / zoom;
29             long double y = ((long double)i - pixel_shift_y) / zoom;
30             complex_number c;
31             c.real = x;
32             c.imaginary = y;
33             complex_number z = c;
34             int iterations = 0; //keep track of the number of iterations
35             for (int k = 0; k < precision; k++)
36             {
37                 complex_number z2;
38                 z2.real = z.real * z.real - z.imaginary * z.imaginary;
39                 z2.imaginary = 2 * z.real * z.imaginary;
40                 z2.real += c.real;
41                 z2.imaginary += c.imaginary;
42                 z = z2;
43                 iterations++;
44                 if (z.real * z.real + z.imaginary * z.imaginary > 4)
45                     break;
46             }
47             //color pixel based on the number of iterations
48             if (iterations < precision / 4.0f)
49             {
50                 vertexarray[i*width + j].position = sf::Vector2f(j, i);
51                 sf::Color color(iterations * 255.0f / (precision / 4.0f), 0, 0);
52                 vertexarray[i*width + j].color = color;
53             }
54             else if (iterations < precision / 2.0f)
55             {
56                 vertexarray[i*width + j].position = sf::Vector2f(j, i);
57                 sf::Color color(0, iterations * 255.0f / (precision / 2.0f), 0);
58                 vertexarray[i*width + j].color = color;
59             }
60             else if (iterations < precision)
61             {
62                 vertexarray[i*width + j].position = sf::Vector2f(j, i);

```

```

59     }
60     else if (iterations < precision)
61     {
62         vertexarray[i*width + j].position = sf::Vector2f(j, i);
63         sf::Color color(0, 0, iterations * 255.0f / precision);
64         vertexarray[i*width + j].color = color;
65     }
66     }
67 }
68 }
69
70 int main()
71 {
72     sf::String title_string = "Mandelbrot Set Plotter";
73     sf::RenderWindow window(sf::VideoMode(width, height), title_string);
74     window.setFramerateLimit(30);
75     sf::VertexArray pointmap(sf::Points, width * height);
76
77     float zoom = 300.0f;
78     int precision = 100;
79     int x_shift = width / 2;
80     int y_shift = height / 2;
81
82     _int64 ctr1 = 0, ctr2 = 0, freq = 0;
83     if (QueryPerformanceCounter((LARGE_INTEGER *)&ctr1) != 0)
84     {
85
86         QueryPerformanceCounter((LARGE_INTEGER *)&ctr2);
87         std::cout << "Start " << ctr1 << std::endl;
88         std::cout << "End " << ctr2 << std::endl;
89         QueryPerformanceFrequency((LARGE_INTEGER *)&freq);
90         std::cout << "ctr1 - ctr2 = " << ctr1 - ctr2 << endl;
91         std::cout << "QueryPerformanceCounter minimum resolution : 1/ " << freq << "seconds" << std::endl;
92         std::cout << "Function takes time: " << ((ctr2 - ctr1) * 1.0 / freq) * 1000000 << " Microseconds." << std::endl;
93         std::cout << endl;
94
95         //generate_mandelbrot_set(pointmap, x_shift, y_shift, precision);
96         generate_mandelbrot_set(pointmap, x_shift, y_shift, precision, zoom);
97
98         while (window.isOpen())
99         {
100             sf::Event event;
101             while (window.pollEvent(event))
102             {
103                 if (event.type == sf::Event::Closed)
104                     window.close();
105             }
106
107             //zoom into area that is left clicked
108             if (sf::Mouse::isButtonPressed(sf::Mouse::Left))
109             {
110                 sf::Vector2i position = sf::Mouse::getPosition(window);
111                 x_shift -= position.x - x_shift;
112                 y_shift -= position.y - y_shift;
113                 zoom *= 2;
114                 precision += 200;
115                 #pragma omp parallel for
116                 for (int i = 0; i < width*height; i++)
117                 {
118                     pointmap[i].color = sf::Color::Black;
119                 }
120                 generate_mandelbrot_set(pointmap, x_shift, y_shift, precision, zoom);


```

```

117             {
118                 pointmap[i].color = sf::Color::Black;
119             }
120             generate_mandelbrot_set(pointmap, x_shift, y_shift, precision, zoom);
121         }
122         window.clear();
123         window.draw(pointmap);
124         window.display();
125     }
126 }
127 else
128 {
129     DWORD dwError = GetLastError();
130     std::cout << "Error vaflue = " << dwError << std::endl;
131 }
132 system("PAUSE");
133
134 return 0;
135 }
136

```

Above 3 pictures are C++ code for Mandelbrot. I have used visual studio to create an image and write the code for Mandelbrot set. To produce a Mandelbrot image I used SFML library in the visual studio. Then we were to measure the Query performance of the Mandelbrot set. Below are pictures of Mandelbrot set image and its query performances;



Query Performances:

```
Start 94385839965
End 94385839983
ctr1 - ctr2 = -18
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 7.08924 Microseconds.
```

The query performance of the above picture is 7.08924 microseconds. This query performance was done when Qpar and vectorization were disabled.

```
Start 95043595227
End 95043595247
ctrl - ctr2 = -20
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 7.87693 Microseconds.
```

The query performance of the above picture is 7.87693 microseconds. This query performance was done when Qpar was No and for vectorization I used SSE2. This performance took a little more time since the Mandelbrot set needed to be vectorized.

```
Start 96120983720
End 96120983740
ctrl - ctr2 = -20
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 7.87693 Microseconds.
```

The query performance of the above picture is 7.87693 microseconds. This query performance was done when Qpar was enable and for vectorization I used AVX2. It has same query performance as the above picture.

```
Start 96373024821
End 96373024844
ctrl - ctr2 = -23
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 9.05847 Microseconds.
```

The query performance of the above picture is 9.05847 microseconds. This query performance was done when Qpar was enable and for vectorization I used AVX. This performance took the most time for the Mandelbrot set to be vectorized.

Mandelbrot AVX Intrinsic:

To make the speed faster of the Mandelbrot set, we can use AVX intrinsic instructions. Below is the function for the AVX intrinsic instruction for Mandelbrot set.

```
void MandelbrotAVX(float x1, float y1, float x2, float y2, int width, int height, int maxIters, unsigned short * image)
```

```
{
```

```
    float dx = (x2-x1)/width;
```

```
    float dy = (y2-y1)/height;
```

```
    // round up width to next multiple of 8
```

```
    int roundedWidth = (width+7) & ~7UL;
```

```
    float constants[] = {dx, dy, x1, y1, 1.0f, 4.0f};
```

```
    __m256 ymm0 = _mm256_broadcast_ss(constants); // all dx
```

```
    __m256 ymm1 = _mm256_broadcast_ss(constants+1); // all dy
```

```
    __m256 ymm2 = _mm256_broadcast_ss(constants+2); // all x1
```

```
    __m256 ymm3 = _mm256_broadcast_ss(constants+3); // all y1
```

```
    __m256 ymm4 = _mm256_broadcast_ss(constants+4); // all 1's (iter increments)
```

```
    __m256 ymm5 = _mm256_broadcast_ss(constants+5); // all 4's (comparisons)
```

```
    float incr[8]={0.0f,1.0f,2.0f,3.0f,4.0f,5.0f,6.0f,7.0f}; // used to reset the i position when j increases
```

```
    __m256 ymm6 = _mm256_xor_ps(ymm0,ymm0); // zero out j counter (ymm0 is just a dummy)
```

```
    for (int j = 0; j < height; j+=1)
```

```
    {
```

```
        __m256 ymm7 = _mm256_load_ps(incr); // i counter set to 0,1,2,...,7
```

```
        for (int i = 0; i < roundedWidth; i+=8)
```

```
        {
```

```
            __m256 ymm8 = _mm256_mul_ps(ymm7, ymm0); // x0 = (i+k)*dx
```

```
            ymm8 = _mm256_add_ps(ymm8, ymm2); // x0 = x1+(i+k)*dx
```

```
            __m256 ymm9 = _mm256_mul_ps(ymm6, ymm1); // y0 = j*dy
```

```
            ymm9 = _mm256_add_ps(ymm9, ymm3); // y0 = y1+j*dy
```


just a dummy)

```
__m256 ymm10 = _mm256_xor_ps(ymm0,ymm0); // zero out iteration counter (ymm0 is
```

```
__m256 ymm11 = ymm10, ymm12 = ymm10;    // set initial xi=0, yi=0
```

```
unsigned int test = 0;
```

```
int iter = 0;
```

```
do
```

```
{
```

```
    __m256 ymm13 = _mm256_mul_ps(ymm11,ymm11); // xi*xi
```

```
    __m256 ymm14 = _mm256_mul_ps(ymm12,ymm12); // yi*yi
```

```
    __m256 ymm15 = _mm256_add_ps(ymm13,ymm14); // xi*xi+yi*yi
```

```
    ymm15 = _mm256_cmp_ps(ymm15,ymm5, _CMP_LT_OQ);    // xi*xi+yi*yi
```

< 4 in each slot

```
    // now ymm15 has all 1s in the non overflowed locations
```

```
    test = _mm256_movemask_ps(ymm15)&255;    // lower 8 bits are comparisons
```

```
    ymm15 = _mm256_and_ps(ymm15,ymm4);        // get 1.0f or 0.0f in each
```

field as counters

```
    ymm10 = _mm256_add_ps(ymm10,ymm15);        // counters for each pixel
```

iteration

```
    ymm15 = _mm256_mul_ps(ymm11,ymm12);    // xi*yi
```

```
    ymm11 = _mm256_sub_ps(ymm13,ymm14);    // xi*xi-yi*yi
```

```
    ymm11 = _mm256_add_ps(ymm11,ymm8);    // xi <- xi*xi-yi*yi+x0 done!
```

```
    ymm12 = _mm256_add_ps(ymm15,ymm15);    // 2*xi*yi
```

```
    ymm12 = _mm256_add_ps(ymm12,ymm9);    // yi <- 2*xi*yi+y0
```

```

        ++iter;

    } while ((test != 0) && (iter < maxIters));

    // convert iterations to output values
    __m256i ymm10i = _mm256_cvtps_epi32(ymm10);

    // write only where needed
    int top = (i+7) < width? 8: width&7;
    for (int k = 0; k < top; ++k)
        image[i+k*j*width] = ymm10i.m256i_i16[2*k];

    // next i position - increment each slot by 8
    ymm7 = _mm256_add_ps(ymm7, ymm5);
    ymm7 = _mm256_add_ps(ymm7, ymm5);

}
ymm6 = _mm256_add_ps(ymm6, ymm4); // increment j counter
}
}

```

After running the AVX intrinsic code, I could observe the better running performance for Mandelbrot set. It was faster than the regular C or C++ code.

Mandelbrot Set with openMP and MPI:

The purpose of this code was to make the performance more efficient and reduce the number of cycles of Mandelbrot set. Since we build a raspberry pi cluster, which has 8 nodes and each node has 4 processes. Which means our code should be 32 times faster using mpi and openMP. Using

MPI we divide the code into 8 pieces. Where each node will be assigned a piece of code to be executed. Since each node has 4 cores. We can write openMP code to divide each node's code into 4 cores. Which means the code that each node has will be divided into 4 parts and assign each part to each core using openMP.

How to Run the Code on the Raspberry Pi Cluster:

1. First We make a directory of our name in the master node.

```
mkdir ~/mpich/{YOUR_NAME}
```

2. Then we send the command to the rest of the nodes

```
parallel-ssh -i -h ~/pssh_hosts mkdir ~/mpich/{YOUR_NAME}
```

3. Then we copy the code from the flash drive into the master code

```
cp /media/pi/{YOUR_FLASH_DRIVE}/{YOUR_CODE.CPP}  
~/mpich/{YOUR_NAME}/
```

4. Then compile the code with mpicc

```
mpicc -o {YOUR_EXEC} {SOURCE.CPP}
```

5. Then copy the executable file to all other nodes

```
parallel-scp -v -h ~/pssh_hosts ~/mpich/{YOUR_NAME}/{YOUR_EXEC}  
~/mpich/{YOUR_NAME}/
```

6. Run the executable file on the master node using mpirun or mpiexec

```
mpirun -n 8 -hostfile ~/host_file ~/mpich/{YOUR_NAME}/{YOUR_EXEC}
```

Mandelbrot in MPI:

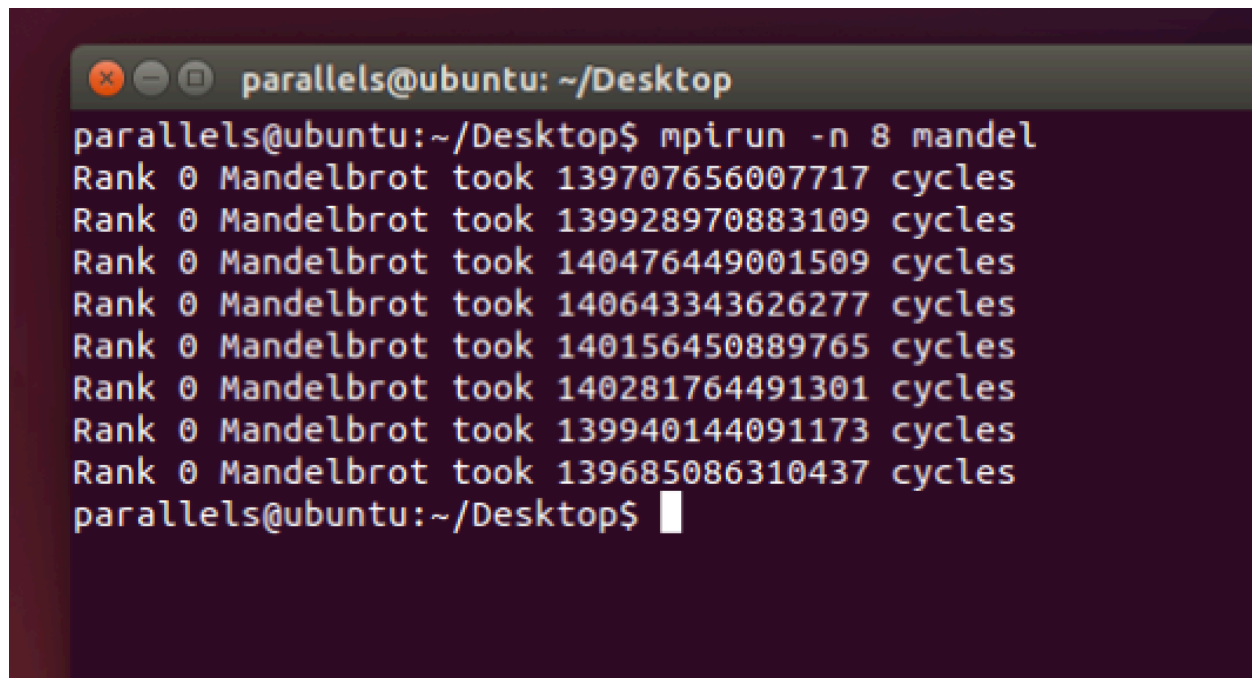
I have 2 ways to run the Mandelbrot code with MPI. First I ran the code using Visual studio compiler and the other one is perf event open in Linux.

```

2 Dir(s) 10,701,007,330 bytes free
Y:\Documents\Visual Studio 2017\Projects\MPITest\Debug>mpiexec -n 8 MPITest.exe
start169942839061
end169943647338
ctr1 - ctr2 = 808277
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 318337 Microseconds.
start169943193801
end169943941927
ctr1 - ctr2 = 748126
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 294647 Microseconds.
start169943108409
end169943966035
ctr1 - ctr2 = 857626
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 337773 Microseconds.
start169943201692
end169944082997
ctr1 - ctr2 = 881305
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 347099 Microseconds.
start169943029671
end169944195749
ctr1 - ctr2 = 1166078
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 459256 Microseconds.
start169943270655
end169944207850
ctr1 - ctr2 = 937195
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 369111 Microseconds.
start169942839025
end169944655009
ctr1 - ctr2 = 1815984
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 715219 Microseconds.
start169943204360
end169944952546
ctr1 - ctr2 = 1748186
QueryPerformanceCounter minimum resolution : 1/ 2539060seconds
Function takes time: 688517 Microseconds.
Press any key to continue . . .

```

As you can see in the picture that every node has different timing. Because it does not have load balance. And I did not assign the specific code to any of the specified node. Since the Mandelbrot is a symmetric image so it takes more time to compute the code from the center of the image. The execution of the top of the image takes less timing than the middle part of the image.

A terminal window titled 'parallels@ubuntu: ~/Desktop' showing the execution of a Mandelbrot set calculation using MPI. The command 'mpirun -n 8 mandel' is entered, followed by eight lines of output, each showing 'Rank 0 Mandelbrot took' followed by a large number of cycles. The numbers range from approximately 139,685,086 to 140,476,490 cycles. The terminal window has a dark background and standard window controls at the top.

```
parallels@ubuntu: ~/Desktop$ mpirun -n 8 mandel
Rank 0 Mandelbrot took 139707656007717 cycles
Rank 0 Mandelbrot took 139928970883109 cycles
Rank 0 Mandelbrot took 140476449001509 cycles
Rank 0 Mandelbrot took 140643343626277 cycles
Rank 0 Mandelbrot took 140156450889765 cycles
Rank 0 Mandelbrot took 140281764491301 cycles
Rank 0 Mandelbrot took 139940144091173 cycles
Rank 0 Mandelbrot took 139685086310437 cycles
parallels@ubuntu:~/Desktop$
```

This image is from Linux terminal. Where we used perf event open command to perform the timing. Since this Mandelbrot code has load balancing, where each processes is assigned different parts of the code. What load balance does is it divides the code into 16 pieces, then assign the middle and upper part of the code to the process; Since upper part takes less than middle part of the image in the Mandelbrot set.

Mandelbrot in MPI and openMP:

This Mandelbrot set code is written using MPI and openMP. The purpose was to reduce the numbers of cycles for the Mandelbrot set and improve the speed and efficiency. To record the timing and number of cycles, I used perf event open in Linux. Which keeps track of the timing of the execution of the code.

```

1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <sys/ioctl.h>
6  #include <linux/perf_event.h>
7  #include <asm/unistd.h>
8  #include <stdlib.h>
9  #include <math.h>
10 #include <omp.h>
11 #include <setjmp.h>
12 #include <tiffio.h>
13 #include <stdbool.h>
14 #include <float.h>
15 #include <mpi.h>
16
17 static long
18 perf_event_open(struct perf_event_attr *hw_event, pid_t pid,
19                int cpu, int group_fd, unsigned long flags)
20 {
21     int ret;
22
23     ret = syscall(__NR_perf_event_open, hw_event, pid, cpu,
24                 group_fd, flags);
25     return ret;
26 }
27
28
29 #define nx 6000 //Resolution in the X dimention
30 #define ny 6000 //Resolution in the Y dimention
31
32 int maxiter= 1000; //max number of iterations to test for an escaping point
33 int myRank;
34 int commSize;
35
36 void calc_pixel_value(int calcny, int calcnx, int calcMSet[calcnx*calcny], int calcmaxiter);
37 void calcSet(int startIdx, int endIdx, int chunkSize);
38
39
40 int main(int argc, char **argv)
41 {
42     struct perf_event_attr pe;
43     long long count;
44     int fd;
45
46     memset(&pe, 0, sizeof(struct perf_event_attr));
47     pe.type = PERF_TYPE_HARDWARE;
48     pe.size = sizeof(struct perf_event_attr);
49     pe.config = PERF_COUNT_HW_INSTRUCTIONS;
50     pe.disabled = 1;
51     pe.exclude_kernel = 1;
52     pe.exclude_hv = 1;
53
54     fd = perf_event_open(&pe, 0, -1, -1, 0);
55     if (fd == -1) {
56         fprintf(stderr, "Error opening leader %llx\n", pe.config);
57         exit(EXIT_FAILURE);
58     }
59
60     ioctl(fd, PERF_EVENT_IOC_RESET, 0);
61     ioctl(fd, PERF_EVENT_IOC_ENABLE, 0);
62

```

```

64 MPI_Init(&argc, &argv);
65 //Here we create a 1 dimensional array that will be the total size of the image (in pixels). Because
66 //this array can get so big, we need to declare it on the heap and NOT the stack (i.e. how one would normally
67 //declare an array. So we use malloc();
68 int *MSet = (int*)malloc(nx*ny*sizeof(int));
69 memset(MSet, 0, nx*ny*sizeof(int));
70 //Get rank number and total number of processes that were launched
71 MPI_Comm_size(MPI_COMM_WORLD, &commSize);
72 MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
73 //Check to see if the resolution and number of processes are compatible.
74 if (myRank == 0)
75 {
76     if ((nx*ny)%commSize != 0)
77     {
78         printf("Incompatible number of processes requested\nExiting...\n");
79         exit(EXIT_FAILURE);
80     }
81 }
82 int totalRes = (nx*ny);
83 int chunkSize = totalRes/(commSize-1);
84
85 if (myRank == 0)
86     printf("Starting Calculation of Mandelbrot set...\n");
87
88 int i;
89 MPI_Barrier(MPI_COMM_WORLD);
90 //This only happens on the manager rank (rank 0). Basically it will compute the boundaries of the section of MSet
91 //compute nodes will process and the size of the region and send them to the appropriate process. Then it will
92 //the completed section of the array it was assigned. When its received it will copy it into the right place in
93 if (myRank == 0)
94 {
95     int myStart = 0, myEnd = 0;
96     int* tempMSet = (int*)malloc((chunkSize)*sizeof(int));
97     for (i = 1; i<commSize; i++)
98     {
99         myStart = (chunkSize/ny)*(i-1);
100         myEnd = (myStart + (chunkSize/ny));
101
102         //send to each node
103         //tag 0 = startIdx
104         //tag 1 = endIdx
105         //tag 2 = chunkSize
106
107         MPI_Send(&myStart, 1, MPI_INT, i, 0, MPI_COMM_WORLD);
108         MPI_Send(&myEnd, 1, MPI_INT, i, 1, MPI_COMM_WORLD);
109         MPI_Send(&chunkSize, 1, MPI_INT, i, 2, MPI_COMM_WORLD);
110         printf("Sent chunk data to node %d\n", i);
111     }
112     int sender = -1;
113     MPI_Status status;
114     //Receive the completed chunks and memcpy it onto MSet
115     for (i = 1; i<commSize; i++)
116     {
117         MPI_Recv(tempMSet, chunkSize, MPI_INT, MPI_ANY_SOURCE, 3, MPI_COMM_WORLD, &status);
118         sender = status.MPI_SOURCE;
119         printf("Received chunk from %d\n", sender);
120         memcpy(MSet+(chunkSize*(sender-1)), tempMSet, chunkSize*sizeof(int));
121         printf("Memcpy'd data from rank %d into MSet\n", sender);
122         sender = -1;
123     }

```

```

126 //set of points then send the completed region back to the manager rank.
127 else if (myRank != 0)
128 {
129     int myStart = 0;
130     int myEnd = 0;
131     int chunkSize = 0;
132     MPI_Status status;
133     MPI_Recv(&myStart, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
134     MPI_Recv(&myEnd, 1, MPI_INT, 0, 1, MPI_COMM_WORLD, &status);
135     MPI_Recv(&chunkSize, 1, MPI_INT, 0, 2, MPI_COMM_WORLD, &status);
136     //DO THE MANDEL
137     calcSet(myStart, myEnd, chunkSize);
138 }
139 MPI_Barrier(MPI_COMM_WORLD);
140 //Once the final array (MSet) is completed, it will call the same tiff writing function from tiff.c
141 if (myRank == 0)
142 {
143     printf("Starting to write image\n");
144     calc_pixel_value(nx,ny,MSet,maxiter);
145 }
146 //End the MPI instance
147 MPI_Finalize();
148 }
149
150 void calcSet(int startIdx, int endIdx, int chunkSize)
151 {
152     int position = 0;
153     printf("Calculating Mandelbrot chunk on process %d\n", myRank);
154     //Here, you can change the minimum and maximum values for the "window" that the image sees. increa
155     //and you will zoom out. Decrease them, and you will effectively zoom into the image.
156     int xmin=-3, xmax= 1;
157     int ymin=-2, ymax= 2;
158     double threshold = 1.0;
159     double dist = 0.0;
160     int ix, iy;
161     double cx, cy;
162     int iter, i = 0;
163     double x,y,x2,y2 = 0.0;
164     double temp=0.0;
165     double xder=0.0;
166     double yder=0.0;
167     double huge = 100000;
168     bool flag = false;
169     const double overflow = DBL_MAX;
170     double delta = (threshold*(xmax-xmin))/(double)(nx-1);
171     int rowSize = ny/(commSize-1);
172     int *localMSet = (int*)malloc((chunkSize)*sizeof(int));
173     int count = 0;
174     #pragma omp parallel
175     {
176         double xorbit[maxiter+1];
177         xorbit[0] = 0.0;
178         double yorbit[maxiter+1];
179         yorbit[0] = 0.0;
180         //Start OpenMP code
181         //We use a nested loop here to effectively traverse over each part of the grid (pixel of the ima
182         //determined and then used as the basis of the computaion. Effectively, it will loop over each p
183         //the value that the mathematical function returns on each iteration it will determine whether o
184         //number.) or not. If it takes few iterations to escape then it will decide that this point is N
185         //index in MSet. If it takes nearly all or all of the iterations to escape, then it will decide
186         //put a 1 in its place in MSet.
187         //The use of the OpenMP pragma here will divide up the iterations between threads and execute th

```



```

189     #pragma omp for
190     for (iy = startIdx; iy<endIdx; iy++)
191     {
192         cy = ymin+iy*(ymax-ymin)/(double)(ny-1);
193         for (ix = 0; ix<=(nx-1); ix++)
194         {
195             iter = 0;
196             i = 0;
197             x = 0.0;
198             y = 0.0;
199             x2 = 0.0;
200             y2 = 0.0;
201             temp = 0.0;
202             xder = 0.0;
203             yder = 0.0;
204             dist = 0.0;
205             cx = xmin +ix*(xmax-xmin)/(double)(ny-1);
206             //This is the main loop that determines whether or not the point escapes or not. It breaks out
207             for (iter =0; iter<maxiter; iter++)
208             {
209                 temp = x2-y2 +cx;
210                 y = 2.0*x*y+cy;
211                 x = temp;
212                 x2 = x*x;
213                 y2 = y*y;
214                 xorbit[iter+1]=x;
215                 yorbit[iter+1]=y;
216                 if (x2+y2>huge) break; //if point escapes then break to next loop
217             }
218             //if the point escapes, find the distance from the set, just incase its close to the set. if
219             if (x2+y2>=huge)
220             {
221                 xder, yder = 0;
222                 i = 0;
223                 flag = false;
224
225                 for (i=0;i<=iter && flag==false;i++)
226                 {
227                     temp = 2.0*(xorbit[i]*xder-yorbit[i]*yder)+1;
228                     yder = 2.0*(yorbit[i]*xder+xorbit[i]*yder);
229                     xder = temp;
230                     flag = fmax(fabs(xder), fabs(yder)) > overflow;
231                 }
232                 if (flag == false)
233                 {
234                     dist=(log(x2+y2)*sqrt(x2+y2))/sqrt(xder*xder+yder*yder);
235                 }
236             }
237
238             //Assign the appropriate values to MSet in the place relating to the point in question
239             if (dist < delta)
240                 localMSet[count*(nx)+ix] = 1;
241             else
242                 localMSet[count*(nx)+ix] = 0;
243
244         }
245         count++;
246     }
247 }
248 printf("Sending calculated set back to master from rank %d\n", myRank);
249 //Send the array back to be memcpy'ed into MSet

```

```

245         count++;
246     }
247 }
248 printf("Sending calculated set back to master from rank %d\n", myRank);
249 //Send the array back to be memcpy'ed into MSet
250 MPI_Send(localMSet, chunkSize, MPI_INT, 0, 3, MPI_COMM_WORLD);
251
252
253 ioctl(fd, PERF_EVENT_IOC_DISABLE, 0);
254 read(fd, &count, sizeof(long long));
255
256 printf("Used %lld instructions\n", count);
257
258 close(fd);
259 }
260

```

This is the corresponding for the Mandelbrot code using MPI and openMP. As you can see I have divided the code into the specified node and used openMP to assign the task to all cores in the nodes. This code produces much better results in terms of timing and performance.

Conclusion:

The purpose of the project was write Mandelbrot set code using C++ or C and check the query performance. Then we had to use MPI and openMP to reduce the number of cycles and improve the timing and efficiency of the code. As a result the Mandelbrot code with MPI and openMPI was faster than the rest of the other code.